

190905104

Parth Shukla

1. Write a program in CUDA to multiply two Matrices for the following specifications:

- a. Each row of resultant matrix to be computed by one thread.
- b. Each column of resultant matrix to be computed by one thread.
- c. Each element of resultant matrix to be computed by one thread.
- d. Perform matrix multiplication using 2D Grid and 2D Block.

```
#include <cuda.h>
```

```
#include <stdlib.h>
```

```
#include <stdio.h>
```

```
__global__ void MatrixMulKernel_a(int *d_M, int *d_N, int *d_P, int wa, int wb)
```

```
{
```

```
    int ridA = threadIdx.x;
```

```
    int sum;
```

```
    for (int cidB = 0; cidB < wb; cidB++)
```

```
    {
```

```
        sum = 0;
```

```
        for (int k = 0; k < wa; k++)
```

```
        {
```

```
            sum += d_M[ridA * wa + k] * d_N[k * wb + cidB];
```

```
        }
```

```
        d_P[ridA * wb + cidB] = sum;
```

```
    }
```

```
}
```

```
__global__ void MatrixMulKernel_b(int *d_M, int *d_N, int *d_P, int ha, int wa)
```

```
{
```

```
    int cidB = threadIdx.x;
```

```
    int wb = blockDim.x;
```

```

int sum;
for (int ridA = 0; ridA < ha; ridA++)
{
    sum = 0;
    for (int k = 0; k < wa; k++)
    {
        sum += d_M[ridA * wa + k] * d_N[k * wb + cidB];
    }
    d_P[ridA * wb + cidB] = sum;
}
}

```

```

__global__ void
MatrixMulKernel_c(int *d_M, int *d_N, int *d_P, int Width)
{

    int col = threadIdx.x;
    int row = threadIdx.y;
    int k = 0;
    for (int i = 0; i < width; i++)
    {
        k += A[row * width + i] * B[col + i * width];
    }
    C[row * width + col] = k;
}

```

```

__global__ void MatrixMulKernel_d(const int *a, const int *b, int *c, int m, int n, int
o)
{
    // row and col calculations
    int row = blockIdx.y * blockDim.y + threadIdx.y;
    int col = blockIdx.x * blockDim.x + threadIdx.x;
    c[row * o + col] = 0;
}

```

```

// calculating one element
for (int k = 0; k < n; k++)
{
    c[row * o + col] += a[row * n + k] * b[k * o + col];
}
}

```

```

__host__ void clearMatrix(int *A, int width)
{
    for (int i = 0; i < width; i++)
    {
        for (int j = 0; j < width; j++)
        {
            A[i * width + j] = 0;
        }
    }
}

```

```

__host__ __device__ void printMatrix(const char *string, int *A, int width)
{
    printf("%s\n", string);
    for (int i = 0; i < width; i++)
    {
        for (int j = 0; j < width; j++)
        {
            printf("%d, ", A[i * width + j]);
        }
        printf("\n");
    }
    printf("\n");
}

```

```

void multiplyMatrix(int *h_A, int *h_B, int *h_C, int width)

```

```

{
    int *d_A, *d_B, *d_C;
    int size = width * width * sizeof(int);
    cudaMalloc((void **)&d_A, size);
    cudaMalloc((void **)&d_B, size);
    cudaMalloc((void **)&d_C, size);
    cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_B, h_B, size, cudaMemcpyHostToDevice);
    cudaMemcpy(d_C, h_C, size, cudaMemcpyHostToDevice);
    dim3 dimBlock(1, 1, 1);
    dim3 dimGrid(1, 1, 1);
    dimBlock.x = 1;
    dimBlock.y = width;
    dimBlock.z = 1;
    MatrixMulKernel_a<<<1, width>>>(d_A, d_B, d_C, width, width);
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
    printMatrix("A*B: (from a kernel): ", h_C, width);
    clearMatrix(h_C, width);
    cudaMemcpy(d_C, h_C, size, cudaMemcpyHostToDevice);
    dimBlock.x = width;
    dimBlock.y = 1;
    dimBlock.z = 1;
    MatrixMulKernel_b<<<1, width>>>(d_A, d_B, d_C, width, width);
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
    printMatrix("A*B: (from b kernel): ", h_C, width);
    clearMatrix(h_C, width);
    cudaMemcpy(d_C, h_C, size, cudaMemcpyHostToDevice);
    dimBlock.x = width;
    dimBlock.y = width;
    dimBlock.z = 1;
    MatrixMulKernel_c<<<(1, 1), (width, width)>>>(d_A, d_B, d_C, width);
    cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);

```

```

printMatrix("A*B: (from c kernel): ", h_C, width);
clearMatrix(h_C, width);
cudaMemcpy(d_C, h_C, size, cudaMemcpyHostToDevice);
dimBlock.x = width;
dimBlock.y = width;
dimBlock.z = 1;
MatrixMulKernel_d<<<dimBlock, dimGrid>>>(d_A, d_B, d_C, width, width,
width);
cudaMemcpy(h_C, d_C, size, cudaMemcpyDeviceToHost);
printMatrix("A*B: (from d kernel): ", h_C, width);
clearMatrix(h_C, width);
cudaFree(d_A);
cudaFree(d_B);
cudaFree(d_C);
}
int main()
{
    int *A, *B, *C;
    int width = 3;
    int size = width * width * sizeof(int);
    A = (int *)calloc(width * width, sizeof(int));
    B = (int *)calloc(width * width, sizeof(int));
    C = (int *)calloc(width * width, sizeof(int));
    int k = 1;
    for (int i = 0; i < width; i++)
    {
        for (int j = 0; j < width; j++)
        {
            A[i * width + j] = rand() % 10;
            B[i * width + j] = rand() % 11;
            k++;
        }
    }
}

```

```

printMatrix("A:", A, width);
printMatrix("B:", B, width);
multiplyMatrix(A, B, C, width);
return 0;
}

```

```

student@dmlab-hp-280-10: ~/190905104_ParthShukla_PCAP/...
student@dmlab-hp-280-10: ~/1909051... x student@dmlab-hp-280-10: ~/1909051... x
D-190905104@lplab-ProLiant-DL380-G6:~/week6$ ./a.out
A:
3, 7, 3,
6, 9, 2,
0, 3, 0,

B:
10, 2, 4,
6, 1, 7,
3, 4, 10,

A*B: (from a kernel):
81, 25, 91,
120, 29, 107,
18, 3, 21,

A*B: (from b kernel):
81, 25, 91,
120, 29, 107,
18, 3, 21,

A*B: (from c kernel):
81, 25, 91,
120, 29, 107,
18, 3, 21,

```

2. Write a program in CUDA to read a sentence with equal length words. Count the number of times a given word is repeated in this sentence. (Use Atomic function).

Sample string: Pcap EEFM exam Pcap test Pcap

Word: Pcap

Given word repeated 3 times

```

#include <cuda.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
__global__ void word_count_kernel(char *str, char *key, int *word_indices, int *result)

```

```

{
    int idx = threadIdx.x + blockIdx.x * blockDim.x;
    // get idx'th word
    int si = word_indices[idx];
    int ei = word_indices[idx + 1];
    char word[100];
    int i = 0;
    for (i = 0; i < (ei - si - 1); i++)
    {
        word[i] = str[si + 1 + i];
    }
    word[i] = '\0';
    // compare word and key
    int i1 = 0;
    int i2 = 0;
    int is_equal = 1;
    while (word[i1] != '\0' && key[i2] != '\0')
    {
        if (word[i1] == key[i2])
        {
            i1++;
            i2++;
        }
        else
        {
            is_equal = 0;
            break;
        }
    }
    if (is_equal == 1)
    {
        atomicAdd(result, 1);
    }
}

int main()
{
    char str[100];
    char key[100];
    printf("Enter input String:\n");
    gets(str);
    printf("Enter a Key:\n");
    gets(key);
    int str_len = strlen(str);
    //   Inputed String : Hello Hi Hello Hi Hello
    //               Key : hi\
// Total occurances of hi\ is 0 D -
//               190905087 @lplab - ProLiant - DL380 - G6 : ~$./ atom Enter input String :

    int key_len = strlen(key);
    int word_count = 0;
    for (int i = 0; i < str_len; i++)
    {

```

```

    if (str[i] == ' ')
    {
        word_count++;
    }
}
word_count--;
int *word_indices = (int *) (malloc(word_count * sizeof(int)));
int wi = -1;
for (int i = 0; i < str_len; i++)
{
    if (str[i] == ' ')
    {
        word_indices[++wi] = i;
    }
}
int result = 0;
char *d_str;
char *d_key;
int *d_word_indices;
int *d_result;
cudaMalloc((void **)&d_str, str_len * sizeof(char));
cudaMalloc((void **)&d_key, key_len * sizeof(char));
cudaMalloc((void **)&d_word_indices, (word_count + 1) * sizeof(int));
cudaMalloc((void **)&d_result, sizeof(int));
cudaMemcpy(d_str, str, str_len * sizeof(char), cudaMemcpyHostToDevice);
cudaMemcpy(d_key, key, key_len * sizeof(char),
           cudaMemcpyHostToDevice);
cudaMemcpy(d_word_indices, word_indices, (word_count + 1) * sizeof(int),
           cudaMemcpyHostToDevice);
cudaMemcpy(d_result, &result, sizeof(int), cudaMemcpyHostToDevice);
word_count_kernel<<<1, word_count>>>(d_str, d_key, d_word_indices,
                                     d_result);
cudaMemcpy(&result, d_result, sizeof(int), cudaMemcpyDeviceToHost);
printf("Inputed String: %s\n", str);
printf("Key: %s\n", key);
printf("Total occurances of %s is %d\n", key, result);
cudaFree(d_str);
cudaFree(d_key);
cudaFree(d_result);
return 0;
}

```

```

D-190905104@lplab-ProLiant-DL380-G6:~/week6$ ./a.out
Enter input String:
Pcap EEFM exam Pcaptest Pcap
Enter a Key:
Pcap
Inputed String: Pcap EEFM exam Pcaptest Pcap
Key: Pcap
Total occurances of Pcap is 2

```